# Week 2 Lab 2 — Sydney AirBnB Exploratory Data Analysis (EDA)

4th November 2026

Ask a traveller to describe their dream Sydney stay, and they'll probably start with the beach view or the walk to the harbour — not the host's response rate or how many other listings the host manages. But this dataset proves that much more shapes a listing's nightly price than bedrooms and a nice photo.

With 34 variables describing every Airbnb listing in Sydney — 20,573 of them in the June 2026 snapshot — the problem is to understand what drives the price of a stay.

Acknowledgments The data comes from Inside Airbnb (https://insideairbnb.com), an independent open-data project that publishes quarterly snapshots of Airbnb's public listings, used here under its Creative Commons Attribution 4.0 licence. This is the Sydney snapshot of 16 June 2026 — the same listings you could find on the site that week.

The explanation of the columns comes from Inside Airbnb's data dictionary.

Exploratory Data Analysis (EDA)

Exploratory Data Analysis (EDA) is an approach used in data science to summarize and visually analyze the characteristics of a data set. It helps to answer important questions related to using machine learning algorithms in a project, selecting the most suitable algorithms, defining feature variables for machine learning, and finding insights that were not evident to stakeholders.

EDA involves multiple exploratory techniques, including :

univariate

bivariate

multivariate visualization

dimensionality reduction

By validating assumptions and identifying patterns, EDA ensures that the data set is ready to be used in a machine learning algorithm. Skipping or doing a shallow work on EDA can lead to inaccurate or inefficient models, choosing the wrong variables, and other negative outcomes. Data profiling is an important aspect of EDA, as it provides a solid understanding of the data set and allows for querying and visualizing the data in various ways.

[1]:
```
# pandas      -> loading and working with the data table
# numpy       -> number types (used when selecting numeric columns)
# matplotlib  -> plotting
# seaborn     -> the heatmaps and pair plot
import pandas as pd
```

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

[2]:
```
# read the Sydney Airbnb listings dataset into a DataFrame called df
df = pd.read_csv('data/sydney_airbnb.csv')
```

[3]:
```
# display the table - 20,573 listings, 34 columns
df
```

[3]:

```
            id            suburb   latitude   longitude  \
0        11156            Sydney -33.867670  151.224970
1        15253            Sydney -33.879640  151.216800
2        44545            Sydney -33.878880  151.214390
3        58506            Mosman -33.817480  151.234840
4      1141803            Sydney -33.887060  151.214030
…        …                 …         …          …
20568  1125557  Sutherland Shire -34.048250  151.156220
20569  1130697       Leichhardt -33.867840  151.171090
20570  1130919         Randwick -33.915460  151.254330
20571  1131054        Woollahra -33.889020  151.236800
20572  1139616       Parramatta -33.776134  151.080627

                     property_type        room_type  accommodates  bathrooms  \
0      Private room in rental unit     Private room             1        NaN
1            Private room in condo     Private room             2        1.0
2                      Entire loft  Entire home/apt             2        NaN
3                Entire guest suite  Entire home/apt             4        1.0
4                      Entire loft  Entire home/apt             2        1.0
…                              …            …           …        …
20568  Private room in rental unit     Private room             2        NaN
20569        Private room in home      Private room             1        NaN
20570                  Entire home  Entire home/apt             6        2.5
20571          Entire rental unit  Entire home/apt             4        1.0
20572        Private room in home      Private room             1        2.0

       bathrooms_text  bedrooms  …  review_scores_cleanliness  \
0       1 shared bath       NaN  …                       4.30
1      1 private bath       NaN  …                       4.54
2              1 bath       1.0  …                       4.86
3              1 bath       1.0  …                       4.36
4              1 bath       1.0  …                       4.87
…                 …       …  …                          …
20568   1 shared bath       NaN  …                       4.70
20569         2 baths       NaN  …                       4.30
20570       2.5 baths       3.0  …                       4.79
20571          1 bath       2.0  …                       4.98
```

```
20572  2 shared baths       NaN  …                        4.90

       review_scores_checkin  review_scores_communication  \
0                       4.80                         4.76
1                       4.82                         4.76
2                       4.92                         4.97
3                       4.90                         4.92
4                       4.96                         4.97
…                          …                            …
20568                   4.89                         4.89
20569                   4.85                         4.81
20570                   5.00                         5.00
20571                   4.98                         4.89
20572                   4.89                         4.90

       review_scores_location  review_scores_value  host_is_superhost  \
0                        4.82                 4.71                  f
1                        4.77                 4.58                  f
2                        4.94                 4.80                  t
3                        4.77                 4.73                  t
4                        5.00                 4.83                  t
…                           …                    …                  …
20568                    4.87                 4.63                  f
20569                    4.78                 4.56                  f
20570                    5.00                 4.79                  f
20571                    4.97                 4.84                  f
20572                    4.87                 4.87                  f

       host_identity_verified  hosts_time_as_host_years  host_listings_count  \
0                           t                        15                    1
1                           t                        14                    1
2                           t                        15                    1
3                           t                        15                    1
4                           t                        13                    1
…                           …                         …                    …
20568                       t                        11                    1
20569                       t                        13                    2
20570                       t                        10                    1
20571                       t                        13                    2
20572                       t                        13                    2

       calculated_host_listings_count
0                                   1
1                                   1
2                                   1
3                                   1
4                                   1
```

```
…                                   …
20568                               1
20569                               2
20570                               1
20571                               2
20572                               2

[20573 rows x 34 columns]
```

A good first step is a basic overview of the data — a quick, informal look at its shape, ranges and averages. The describe() method does this in one line.

[4]:
```
# summary statistics for every numeric column:
# count, mean, spread (std), smallest, largest, and the quartiles
df.describe()
```

[4]:
```
                 id      latitude     longitude  accommodates     bathrooms  \
count  2.057300e+04  20573.000000  20573.000000  20573.000000  14849.000000
mean   9.671166e+17    -33.853299    151.170113      4.036358      1.505859
std    6.240037e+17      0.088507      0.118498      2.510689      0.809489
min    1.115600e+04    -34.110210    150.630490      1.000000      0.500000
25%    5.414616e+07    -33.896820    151.122470      2.000000      1.000000
50%    1.166776e+18    -33.875670    151.204370      4.000000      1.000000
75%    1.522393e+18    -33.809820    151.248915      5.000000      2.000000
max    1.708696e+18    -33.381266    151.340140     16.000000     21.000000

           bedrooms          beds         price  minimum_nights  \
count  17466.000000  15529.000000  17786.000000    20565.000000
mean       2.043055      2.362483    431.646596       15.444396
std        1.192822      1.702624    893.648969       34.439280
min        0.000000      1.000000      1.740000        1.000000
25%        1.000000      1.000000    169.725000        1.000000
50%        2.000000      2.000000    291.000000        2.000000
75%        3.000000      3.000000    445.330000        4.000000
max       14.000000     50.000000  52850.000000      500.000000

       availability_30  …  review_scores_rating  review_scores_accuracy  \
count     20573.000000  …          17247.000000            17245.000000
mean         15.102562  …              4.713404                4.747029
std          11.328869  …              0.438603                0.416774
min           0.000000  …              1.000000                0.000000
25%           3.000000  …              4.640000                4.680000
50%          16.000000  …              4.830000                4.860000
75%          27.000000  …              4.980000                5.000000
max          30.000000  …              5.000000                5.000000

       review_scores_cleanliness  review_scores_checkin  \
count               17245.000000           17245.000000
```

```
mean                    4.671334              4.811558
std                     0.460948              0.372695
min                     0.000000              0.000000
25%                     4.560000              4.780000
50%                     4.800000              4.920000
75%                     4.970000              5.000000
max                     5.000000              5.000000

       review_scores_communication  review_scores_location  \
count                 17245.000000            17245.000000
mean                      4.833921                4.811273
std                       0.364859                0.332053
min                       0.000000                0.000000
25%                       4.820000                4.770000
50%                       4.950000                4.910000
75%                       5.000000                5.000000
max                       5.000000                5.000000

       review_scores_value  hosts_time_as_host_years  host_listings_count  \
count         17245.000000              20573.000000         20573.000000
mean              4.627076                  5.410149            24.414475
std               0.450956                  4.027088            50.559686
min               0.000000                  0.000000             0.000000
25%               4.510000                  2.000000             1.000000
50%               4.730000                  6.000000             4.000000
75%               4.880000                  9.000000            24.000000
max               5.000000                 15.000000           667.000000

       calculated_host_listings_count
count                    20573.000000
mean                        20.203228
std                         38.468421
min                          1.000000
25%                          1.000000
50%                          3.000000
75%                         20.000000
max                        209.000000

[8 rows x 28 columns]
```

[5]:
```
# just the first 5 rows - a quick look at what the columns hold
df.head()
```

[5]:
```
        id  suburb  latitude  longitude                property_type  \
0    11156  Sydney -33.86767  151.22497  Private room in rental unit
1    15253  Sydney -33.87964  151.21680        Private room in condo
2    44545  Sydney -33.87888  151.21439                  Entire loft
```

```
3    58506  Mosman -33.81748  151.23484          Entire guest suite
4  1141803  Sydney -33.88706  151.21403                 Entire loft

         room_type  accommodates  bathrooms  bathrooms_text  bedrooms  … \
0     Private room             1        NaN   1 shared bath       NaN  …
1     Private room             2        1.0  1 private bath       NaN  …
2  Entire home/apt             2        NaN          1 bath       1.0  …
3  Entire home/apt             4        1.0          1 bath       1.0  …
4  Entire home/apt             2        1.0          1 bath       1.0  …

   review_scores_cleanliness  review_scores_checkin  \
0                       4.30                   4.80
1                       4.54                   4.82
2                       4.86                   4.92
3                       4.36                   4.90
4                       4.87                   4.96

   review_scores_communication  review_scores_location  review_scores_value  \
0                         4.76                    4.82                 4.71
1                         4.76                    4.77                 4.58
2                         4.97                    4.94                 4.80
3                         4.92                    4.77                 4.73
4                         4.97                    5.00                 4.83

   host_is_superhost  host_identity_verified  hosts_time_as_host_years  \
0                  f                       t                        15
1                  f                       t                        14
2                  t                       t                        15
3                  t                       t                        15
4                  t                       t                        13

   host_listings_count  calculated_host_listings_count
0                    1                               1
1                    1                               1
2                    1                               1
3                    1                               1
4                    1                               1

[5 rows x 34 columns]
```

[6]:
```
# (rows, columns) - how big is this dataset?
df.shape
```

[6]:
```
(20573, 34)
```

Missing values

Unlike the first lab's dataset, this one has real gaps. Counting them tells us two different stories:

some values are missing because they cannot exist yet (a listing with no reviews has no review scores), while others were simply never recorded (bathrooms, beds) — an important distinction before deciding how to treat them.

[7]:
```
# count the missing values in every column,
# then keep only the columns that have any, biggest first
missing = df.isnull().sum()
missing[missing > 0].sort_values(ascending=False)
```

[7]:
```
bathrooms                     5724
beds                          5044
review_scores_accuracy        3328
review_scores_cleanliness     3328
review_scores_checkin         3328
review_scores_communication   3328
review_scores_location        3328
review_scores_value           3328
reviews_per_month             3326
review_scores_rating          3326
bedrooms                      3107
price                         2787
estimated_revenue_l365d       2787
bathrooms_text                  25
minimum_nights                   8
dtype: int64
```

Numerical Columns

[8]:
```
# select only the columns that hold numbers (any numeric type)
numeric_features = df.select_dtypes(include=[np.number])

numeric_features.columns
```

[8]:
```
Index(['id', 'latitude', 'longitude', 'accommodates', 'bathrooms', 'bedrooms',
       'beds', 'price', 'minimum_nights', 'availability_30', 'availability_60',
       'availability_90', 'availability_365', 'estimated_occupancy_l365d',
       'estimated_revenue_l365d', 'number_of_reviews', 'number_of_reviews_ltm',
       'reviews_per_month', 'review_scores_rating', 'review_scores_accuracy',
       'review_scores_cleanliness', 'review_scores_checkin',
       'review_scores_communication', 'review_scores_location',
       'review_scores_value', 'hosts_time_as_host_years',
       'host_listings_count', 'calculated_host_listings_count'],
      dtype='object')
```

Categorical Columns

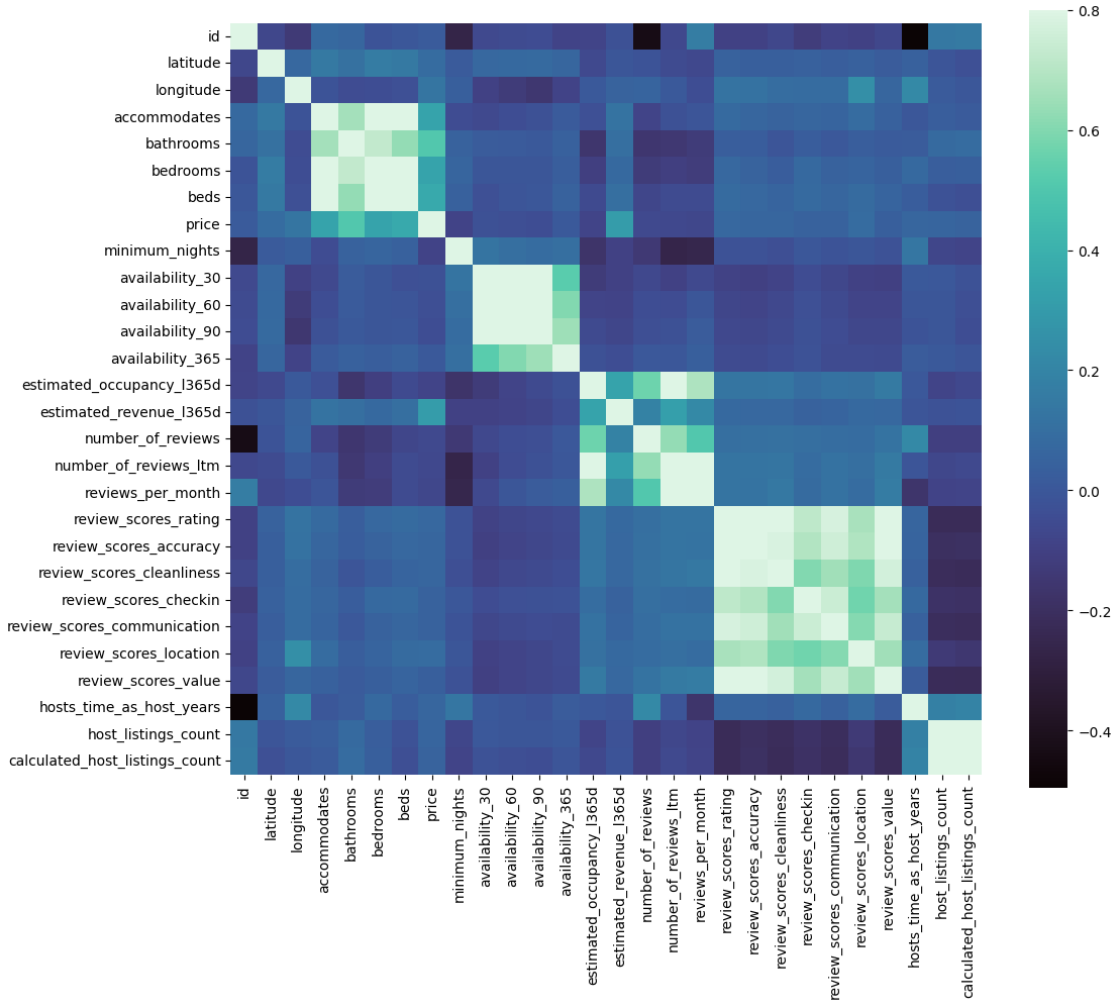
[9]:
```
# select only the columns that hold text - the categorical features
categorical_features = df.select_dtypes(include=['object'])
```

```
categorical_features.columns
```

[9]:
```
Index(['suburb', 'property_type', 'room_type', 'bathrooms_text',
       'host_is_superhost', 'host_identity_verified'],
      dtype='object')
```

[10]:
```
# correlation between every pair of numeric columns, as one big heatmap
# vmax=0.8 caps the colour scale so the strongest links stand out clearly
plt.figure(figsize=(12,10))
sns.heatmap(numeric_features.corr(), square=True, vmax=0.8, cmap='mako')
```

[10]:
```
<Axes: >
```




The heatmap is the best way to get a quick overview of correlated features thanks to seaborn!

The 'availability' columns clearly move together, and so do the size columns ('accommodates', 'bedrooms', 'beds'). These variables give (almost) the same information, and hence it's a case of

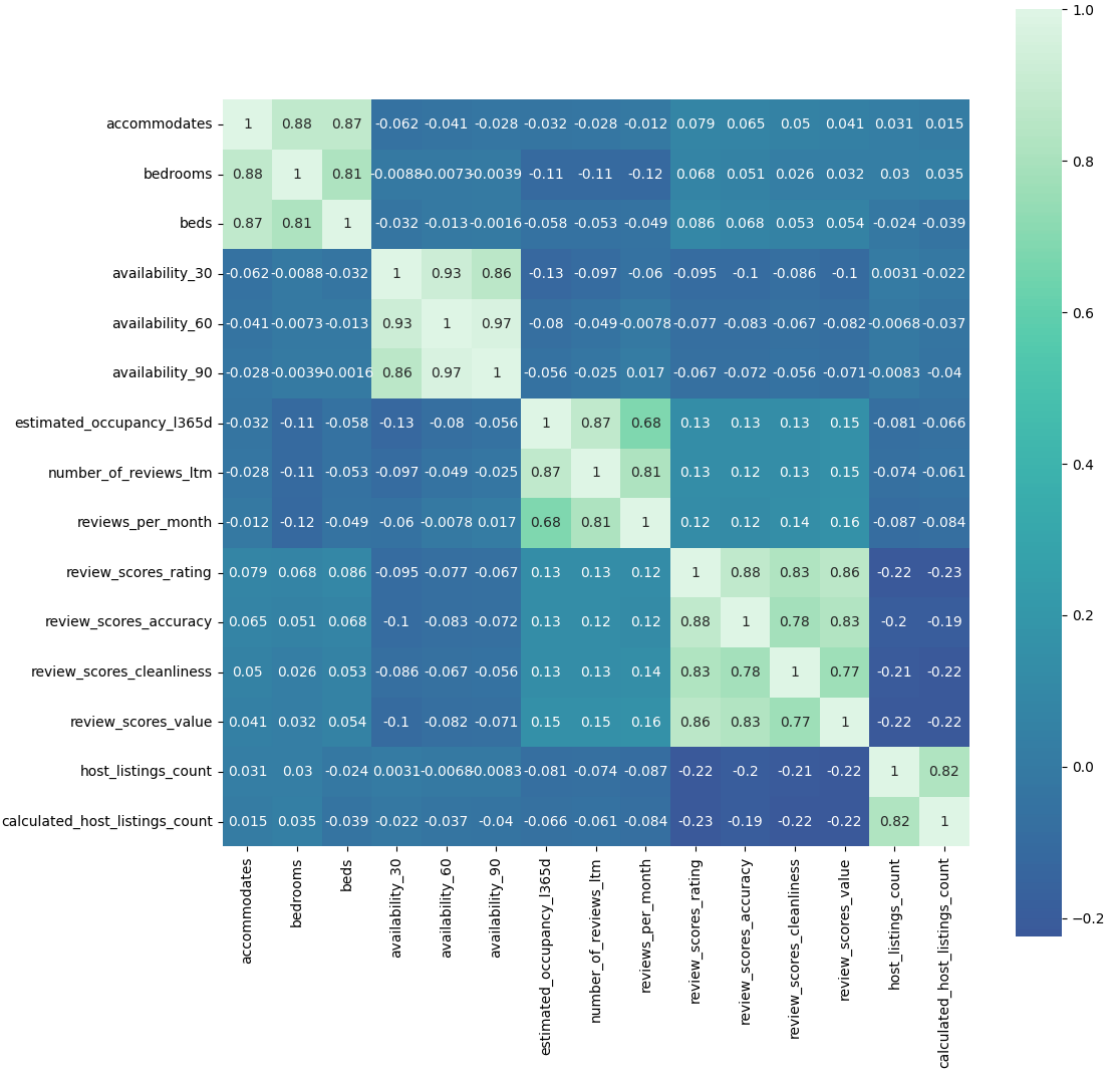
multicollinearity. In ML problems related to feature selection these heatmaps provide an excellent exploratory tool.

[11]:
```
# zoom in on just the strongly related columns
plt.figure(figsize=(12,12))
corr = numeric_features.corr()

# Keep only columns that are strongly correlated (>0.8) with at least one other␣
 ↪column
# (every column counts itself once - its self-correlation is 1.0 - so "> 1"
#  means "correlated with at least one column besides itself")
high_corr_columns = corr.columns[(corr > 0.8).sum() > 1]

# annot=True prints the correlation value inside each square
sns.heatmap(corr.loc[high_corr_columns, high_corr_columns],
            square=True, center=0, annot=True, cmap='mako')
```

[11]: <Axes: >




To understand the code of `corr.columns[(corr > 0.8).sum() > 1]` I will display 4 columns. The first 3 are highly correlated, the last one is not. It can be seen how the first 3 columns have values greater than 0.8 in at least one row, while the last column has no values greater than 0.8 in any other row (excluding with itself).

[12]:
```
display(corr[['accommodates', 'bedrooms', 'beds', 'latitude']].head(10))
```

| | accommodates | bedrooms | beds | latitude |
|---|---|---|---|---|
| id | 0.075259 | -0.017315 | -0.001428 | -0.073978 |
| latitude | 0.143321 | 0.159217 | 0.145853 | 1.000000 |
| longitude | -0.015057 | -0.041684 | -0.033206 | 0.071813 |
| accommodates | 1.000000 | 0.879537 | 0.872992 | 0.143321 |
| bathrooms | 0.660861 | 0.725939 | 0.632816 | 0.108742 |
| bedrooms | 0.879537 | 1.000000 | 0.808719 | 0.159217 |
| beds | 0.872992 | 0.808719 | 1.000000 | 0.145853 |

```
price                0.338285  0.336812  0.364545  0.091884
minimum_nights      -0.046301  0.061232  0.039451  0.011597
availability_30     -0.061819 -0.008801 -0.031669  0.068861
```

[13]:
```
(corr > 0.8).sum()
```

[13]:
```
id                                1
latitude                          1
longitude                         1
accommodates                      3
bathrooms                         1
bedrooms                          3
beds                              3
price                             1
minimum_nights                    1
availability_30                   3
availability_60                   3
availability_90                   3
availability_365                  1
estimated_occupancy_l365d         2
estimated_revenue_l365d           1
number_of_reviews                 1
number_of_reviews_ltm             3
reviews_per_month                 2
review_scores_rating              4
review_scores_accuracy            3
review_scores_cleanliness         2
review_scores_checkin             1
review_scores_communication       1
review_scores_location            1
review_scores_value               3
hosts_time_as_host_years          1
host_listings_count               2
calculated_host_listings_count    2
dtype: int64
```

- 'accommodates', 'bedrooms' and 'beds' are strongly correlated. How many people a listing sleeps is a consequence of how many bedrooms and beds it has — they are like twin brothers (triplets, really). So it is hard to distinguish between them, and we would need only one of these in our analysis.
- 'availability_30', 'availability_60' and 'availability_90' are twins by construction: a night that is free in the next 30 days is also free in the next 60 and 90.
- 'estimated_occupancy_l365d', 'number_of_reviews_ltm' and 'reviews_per_month' move together — the occupancy estimate is calculated from recent reviews, so they are related by definition.
- The review sub-scores ('accuracy', 'cleanliness', 'value') all track the overall 'review_scores_rating': a guest who liked the stay tends to score everything highly.
- 'host_listings_count' and 'calculated_host_listings_count' are two ways of counting the

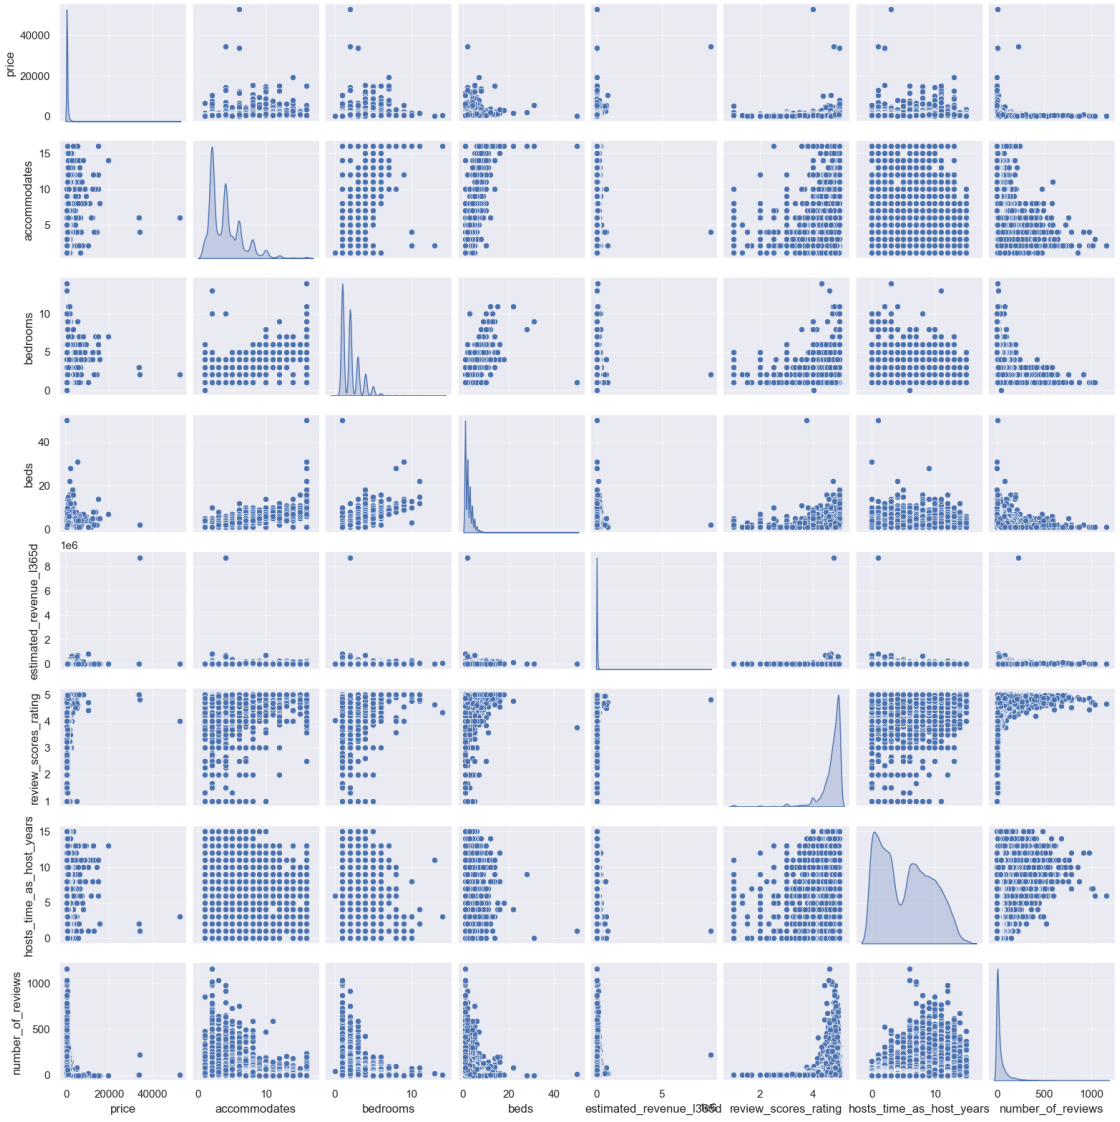
same thing.

Pair Plot

Pair Plots between 'price' and interesting variables Visualisation of 'accommodates','bedrooms','beds','estimated_revenue_l365d','review_scores_rating','hosts_time_as_host_years','number_ features with respect to price in the form of pair plot & scatter pair plot to create better understanding.

[14]:
```
# pair plot: every chosen column plotted against every other
# scatter charts off the diagonal show relationships between pairs;
# the smooth curves on the diagonal show each column's own distribution
# (patience - this one draws over a million points, so it takes a little while)
sns.set_theme()
columns =␣
 ↪['price','accommodates','bedrooms','beds','estimated_revenue_l365d','review_scores_rating',
sns.pairplot(df[columns], height=2, kind='scatter', diag_kind='kde')
```

[14]:
```
<seaborn.axisgrid.PairGrid at 0x15fb5bc20>
```




What did we find?

Reading the price row: size is what moves it — 'accommodates' and 'bedrooms' ( 0.34 each) are the strongest links, with estimated revenue close behind. Just as interesting is what does not move it: 'review_scores_rating' correlates with price at only 0.07 — guests rate cheap stays about as highly as expensive ones. We also found multicollinearity in five separate families — the size trio, the availability windows, the review-activity columns, the review sub-scores, and the two host-listing counts — so a model would need only one column from each family. And several columns have missing values that would need treatment before modelling: listings with no reviews yet have no scores (nothing sensible to impute — the feature does not exist yet), while 'bathrooms' and 'beds' are simply unrecorded for some listings (candidates for imputation).

## 0.1 Extra: plotting the location

Plotly documentation

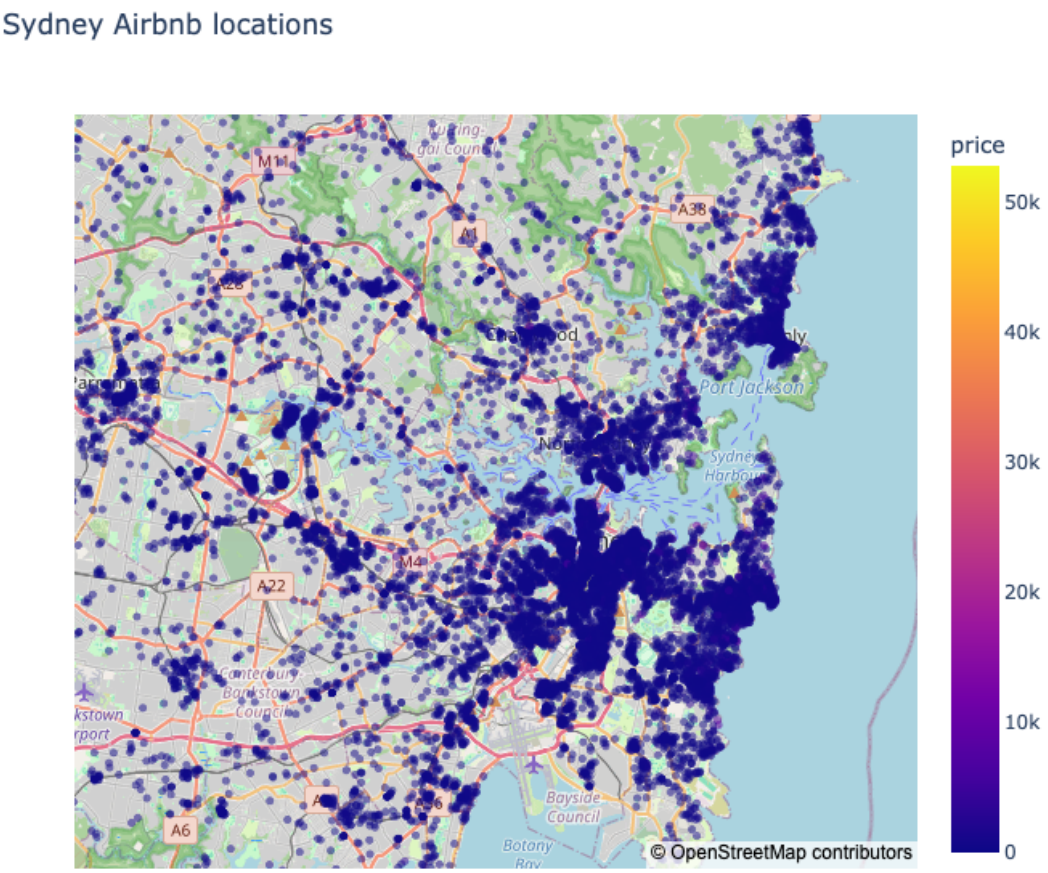
In the map below the outliers really mess up with the price:

[15]:
```
import plotly.express as px

# Exclude listings with missing coordinates
map_df = df.dropna(subset=['latitude', 'longitude', 'price'])

fig = px.scatter_map(
    map_df,
    lat='latitude',
    lon='longitude',
    map_style='open-street-map',
    zoom=10,
    opacity=0.5,
    height=650,
    title='Sydney Airbnb locations',
    color='price',
)

# Keep the interactive notebook map and save a PNG for PDF export.
fig.show(renderer="plotly_mimetype+notebook+png")
```




[16]:
```
from IPython.display import display, Markdown

low_percentile = 10
high_percentile = 90

display(Markdown(
    f"Here the price is grouped into three categories: **low, typical and␣
 ↪high**. "
    f"The low and high groups are defined by percentile **{low_percentile}** "
    f"and percentile **{high_percentile}** of the price distribution. "
    f"The typical group is everything in between."
))
```

Here the price is grouped into three categories: **low, typical and high**. The low and high groups are defined by percentile **10** and percentile **90** of the price distribution. The typical group is

everything in between.

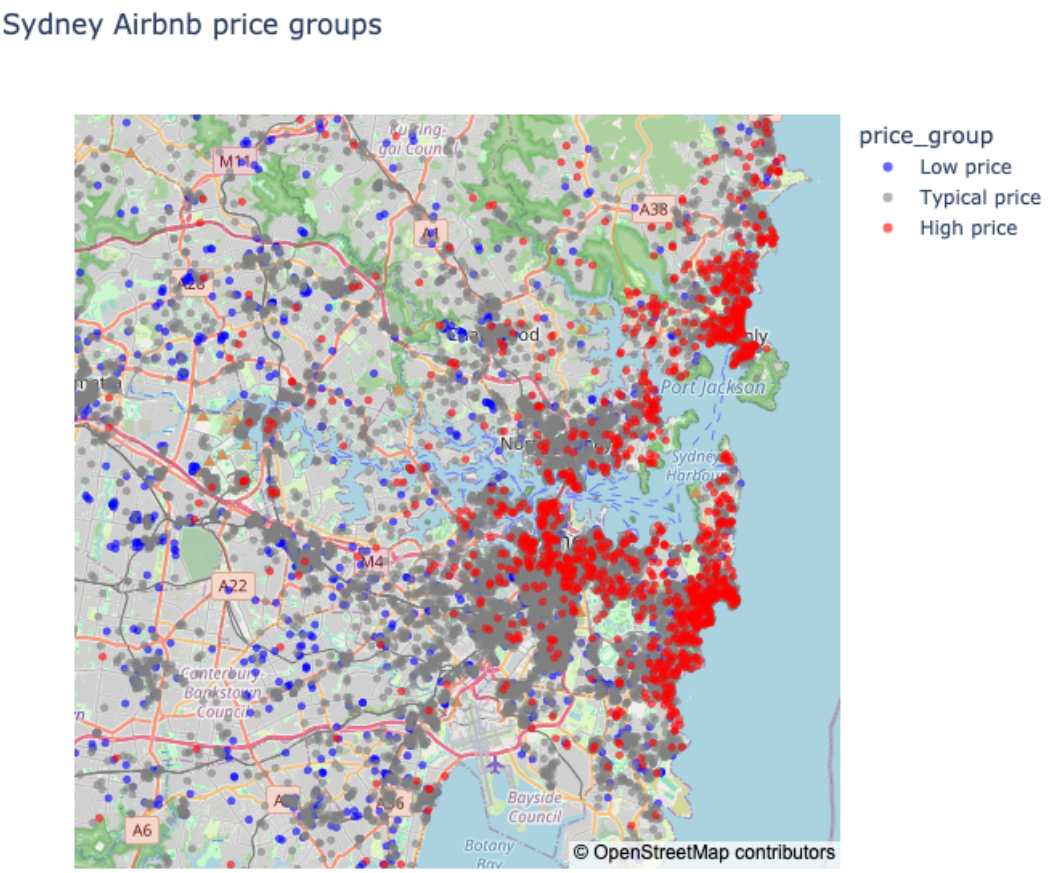
[17]:
```
map_df = map_df.copy()


low = map_df['price'].quantile(low_percentile / 100)
high = map_df['price'].quantile(high_percentile / 100)

map_df['price_group'] = pd.cut(
    map_df['price'],
    bins=[float('-inf'), low, high, float('inf')],
    labels=['Low price', 'Typical price', 'High price'],
)

fig = px.scatter_map(
    map_df,
    lat='latitude',
    lon='longitude',
    color='price_group',
    hover_data=['price'],  # Keep the actual price visible
    category_orders={
        'price_group': ['Low price', 'Typical price', 'High price']
    },
    color_discrete_map={
        'Low price': 'blue',
        'Typical price': 'gray',
        'High price': 'red',
    },
    map_style='open-street-map',
    zoom=10,
    opacity=0.6,
    height=650,
    title='Sydney Airbnb price groups',
)

# Keep the interactive notebook map and save a PNG for PDF export.
fig.show(renderer="plotly_mimetype+notebook+png")
```




[18]:
```
display(Markdown(
    f"The code below caps the prices used for colouring so that extreme values "
    f"don't distort the colour scale. The actual prices are still visible "
    f"in the hover data, but the colour scale is capped at percentile "
    f"**{low_percentile}** and percentile **{high_percentile}**."
))
```

The code below caps the prices used for colouring so that extreme values don't distort the colour scale. The actual prices are still visible in the hover data, but the colour scale is capped at percentile **10** and percentile **90**.

[19]:
```
# Cap only the values used for colouring; preserve actual prices
map_df['price_colour'] = map_df['price'].clip(lower=low, upper=high)

fig = px.scatter_map(
```

```
    map_df,
    lat='latitude',
    lon='longitude',
    color='price_colour',
    range_color=[low, high],
    color_continuous_scale='Viridis',
    hover_data={'price': ':,.2f', 'price_colour': False},
    map_style='open-street-map',
    zoom=10,
    opacity=0.6,
    height=650,
    title='Sydney Airbnb prices',
)

ticks = np.linspace(low, high, 5)
tick_labels = [f'${value:,.0f}' for value in ticks]
tick_labels[0] = f' ${low:,.0f}'
tick_labels[-1] = f' ${high:,.0f}'

fig.update_layout(
    coloraxis_colorbar=dict(
        title='Price',
        tickvals=ticks,
        ticktext=tick_labels,
    )
)

# Keep the interactive notebook map and save a PNG for PDF export.
fig.show(renderer="plotly_mimetype+notebook+png")
```

Sydney Airbnb prices

